

Database and Information Systems

Michal Krátký & Radim Bača

Department of Computer Science
Faculty of Electrical Engineering and Computer Science
VŠB – Technical University of Ostrava

2020/2021



- 1 Concurrency control problems
 - Lost update
 - Uncommitted dependency
 - Inconsistent analysis
- 2 Read/write conflicts
- 3 Concurrency control techniques
- 4 Locking
 - Locking protocol
 - How locking influence concurrency problem
 - Deadlock
- 5 Deadlock



- One criteria is how many users can access the system concurrently:
 - *Single-user* - if only one user can access the database system at one point of time.
 - *Multi-user* - if more users can access the database concurrently.
- We speak about the concurrency if more people access the RDBMS at one point of time.
- Concurrency brings a lot of problems, which influence also the programmers who access the database.



- *Transaction plan* is a sequence of database operations performed during a transaction. We will use these plans to demonstrate problems of concurrency control.
- If we process several plans in one moment, we speak about *concurrent plan* or *parallel plan*.
- To simplify a plan we use the following operations:
 - READ instead of SELECT,
 - WRITE instead of UPDATE, however, these operations can involve also INSERT and DELETE.
- Operations work with tuples (records).

Concurrency control problems



- *Lost update,*
- *Uncommitted dependency,*
- *Inconsistent analysis.*

Lost update



Let us consider these two parallel plans:

Transaction A	Time	Transaction B
READ t	t_1	-
-	t_2	READ t
WRITE t	t_3	-
-	t_4	WRITE t

- 1 Transaction A gets the tuple t in t_1 time.
- 2 Transaction B gets the tuple t in t_2 time.
- 3 Transaction A changes the content of the tuple in t_3 .
- 4 Transaction B rewrites the content of the tuple by its own value in t_4 .

⇒ We lost the update made by the transaction A in t_3 time.

Problem of uncommitted dependency 1/3



- The problem of uncommitted dependency occurs in a case when one transaction reads or updates a tuple which was updated by a transaction which is still not committed.
- Since the transaction is not committed, ROLLBACK can happen.
- In the case of ROLLBACK the first transaction works with values which are not valid.

Problem of uncommitted dependency 2/3



Transaction A	Time	Transaction B
-	t_1	WRITE t
READ t	t_2	-
-	t_3	ROLLBACK

- 1 Transaction A reads an uncommitted value of a tuple t made by the transaction B in time t_2 .
- 2 We ROLLBACK the transaction B in t_3 time.
- 3 The transaction A then works with an invalid tuple t retrieved in time t_2 . The value of t is the value before t_1 .

Problem of uncommitted dependency 3/3



Transaction A	Time	Transaction B
-	t_1	WRITE t
WRITE t	t_2	-
-	t_3	ROLLBACK

- 1 The transaction A writes the tuple t in t_2 time and becomes dependent on an uncommitted change from t_1 time.
- 2 We ROLLBACK the transaction B in t_3 time.
- 3 We lost the change made by the transaction A in t_2 time, since we ROLLBACK to the value of t in time before t_1 .

Problem of inconsistent analysis 1/3



Transactions A and B work with accounts acc_1 , acc_2 and acc_3 . Accounts store values 30, 20, 50 before t_1 .

The transaction A sums the account values, the transaction B moves the value 10 from account acc_1 to account acc_3 .

$acc_1 = 30$	$acc_2 = 20$	$acc_3 = 50$
Transaction A	Time	Transaction B
READ acc_1	t_1	-
$suma = 30$		
READ acc_2	t_2	-
$suma = 50$		
-	t_3	READ acc_3
-	t_4	WRITE $acc_3 = 60$
-	t_5	READ acc_1
-	t_6	WRITE $acc_1 = 20$
-	t_7	COMMIT
READ acc_3	t_8	-
$suma = 110$ but it should be 100		

Problem of inconsistent analysis 2/3



$acc_1 = 30$	$acc_2 = 20$	$acc_1 = 50$
Transaction A	Time	Transaction B
READ acc_1 $suma = 30$	t_1	-
READ acc_2 $suma = 50$	t_2	-
-	t_3	READ acc_3
-	t_4	WRITE $acc_3 = 60$
-	t_5	READ acc_1
-	t_6	WRITE $acc_1 = 20$
-	t_7	COMMIT
READ acc_3 $suma = 110$ but it should be 100	t_8	-

Transaction *A* works with inconsistent database and therefore it makes an inconsistent analysis (sum 110 instead of 100).



Problem of inconsistent analysis 3/3

$acc_1 = 30$	$acc_2 = 20$	$acc_1 = 50$
Transaction A	Time	Transaction B
READ acc_1 $suma = 30$	t_1	-
READ acc_2 $suma = 50$	t_2	-
-	t_3	READ acc_3
-	t_4	WRITE $acc_3 = 60$
-	t_5	READ acc_1
-	t_6	WRITE $acc_1 = 20$
-	t_7	COMMIT
READ acc_3 $suma = 110$ but it should be 100	t_8	-

If we compare it with the problem of uncommitted dependency, we see that the transaction A is not dependent on uncommitted changes of the transaction B since the B committed all changes before A read the acc_3 .

Read/write conflicts



- Concurrency problem occurs when two transactions want to read or write the same tuple.
- There are four different types of conflicts: **RR** (READ-READ), **RW**, **WR**, **WW**.
- In a case of the RR conflict, transactions A and B want to read t . Two read operations can not influence each other, therefore, there is not such a problem.

RW conflict 1/2



- A reads t and B wants to write t . If B proceeds the write, then *problem of inconsistent analysis* can occur.
- The problem of inconsistent analysis is caused by the RW conflict.

Transaction A	Time	Transaction B
READ t	t_1	WRITE t
	t_2	
(another computations)	t_3	

RW conflict 2/2



- If B writes t and A reads t again, then A receives different values.
- We call such a situation *nonrepeatable read* – it is caused by the RW conflict.

Transaction A	Time	Transaction B
READ t	t_1	WRITE t
	t_2	
READ t	t_3	

WR conflict



- A writes t and B wants to read t .
- If B reads t , the **problem of uncommitted dependency** can occur.
- It is called *dirty read*.

Transaction A	Time	Transaction B
WRITE t	t_1	READ t
	t_2	
ROLLBACK ?	t_3	

WW conflict



- A writes t and B wants to write t .
- If B writes t , the **lost update problem** or **problem of uncommitted dependency** can occur (in the case of the transaction B).
- It is called *dirty write*.

Transaction A	Time	Transaction B
WRITE t	t_1	WRITE t
	t_2	
ROLLBACK ?	t_3	

Concurrency control techniques



- There is a number of concurrency control techniques which try to avoid the problems mentioned¹:
 - **locking**,
 - **multi-versioning**,
 - *timestamps*,
 - *validation*.

¹Later we will speak about the serializability of transactions.

Concurrency control techniques



- **Locking** – a pessimistic concurrency control technique: conflicts of parallel transactions are expected.
 - A DBMS manages one copy of data and locks are assigned to transactions.
- **Multi-versioning** – an optimistic concurrency control technique: conflicts of parallel transactions are not expected.
 - A copy of data is created for each transaction, a DBMS should manage which copy is valid.



- **Locking principle:** If a transaction A wants to read or write a tuple in a database, it asks for a lock.
- Another transaction which want to read or write the same tuple has to wait until the transaction A is going to release the lock.

Types of locks



There are two types of locks:

- 1 *Exclusive lock (or write lock)*, we denote X .
- 2 *Shared lock (or read lock)*, we denote S .

There are also other types of locks, however, we consider only these two types to keep it simple now.

Types of locks



The locks are assigned using these rules:

- 1 If a transaction A keeps the **exclusive lock** on a tuple t , then a request of a parallel transaction B on any type of a lock on the same tuple is not processed until A is not going to release the lock. It means, B is **waiting**.
- 2 If a transaction A keeps the **shared lock** on a tuple t , then a request of a parallel transaction B on:
 - 1 **An exclusive lock** is not processed until A is not going to release the lock. It means, B is **waiting**.
 - 2 **A shared lock** is processed immediately. It means, both transactions **keep the shared lock** on t .

Compatibility matrix for a X and S locks



- We can describe these rules using *lock type compatibility matrix*.
- The matrix can be interpreted as follows:
 - Let us consider the tuple t and transaction A which holds the lock corresponding to the first row of the table.
 - A parallel transaction B requests the lock on the tuple t , where the lock type is in the first column of the table.
- Symbol N stands for conflict (tuple t is not locked by the transaction B immediately).

	X	S	-
X	N	N	A
S	N	A	A
-	A	A	A



- Locks are usually requested **implicitly**:
 - The **S** lock is requested automatically before a transaction gets a tuple from a database (for example using the `SELECT` command).
 - The **X** lock is requested before the transaction changes the tuple.
- The change of a tuple can be processed by: `UPDATE`, `INSERT` or `DELETE` operations².

²There are some differences, but we can ignore them now.

Locking protocol 1/3



To solve conflicts of concurrency it is necessary to define the *data access protocol* (or *locking protocol*):

- 1 A transaction which wants to **read a tuple**, has to first request **the shared lock** on this tuple.
- 2 A transaction which wants to **write a tuple**, has to first request **the exclusive lock** on this tuple. If the transaction holds the S lock then this lock is changed to the X lock.



- 3 If the lock requested by the transaction B **can not be granted immediately**:
- The transaction B is switched into **a wait state**.
 - The transaction remains in this state until the request is not processed, i.e. it waits for the release of locks handled by other transactions.
 - A DBMS should prevent the wait state of a transaction forever – this issue is called **livelock** or **starvation**.
 - The most simple way how to solve this issue is to keep **a queue of the requests**: the first transaction asking for the lock first gets it and so on.



- 3** The **exclusive locks** are automatically released at the end of the transaction (finished with COMMIT or ROLLBACK operations). The **shared locks** are usually also released at the end of the transaction.

This protocol is called the **strict two-phase locking**.

Locking – Lost update problem 1/3



Transaction A	Time	Transaction B
READ t (S lock on t assigned)	t_1	-
-	t_2	READ t (S lock on t assigned)
WRITE t (X lock on t requested)	t_3	-
wait	t_4	WRITE t (X lock on t requested)
wait		wait
wait		wait

- The write of a tuple t by the transaction A in time t_3 is not accepted due to its request for the X lock which goes against the S lock of the transaction B .

Locking – Lost update problem 2/3



Transaction A	Time	Transaction B
READ t (S lock on t assigned)	t_1	-
-	t_2	READ t (S lock on t assigned)
WRITE t (X lock on t requested)	t_3	-
wait	t_4	WRITE t (X lock on t requested)
wait		wait
wait		wait

- The transaction A is switched into **the wait state** in t_3 . Due to the same reason also the transaction B is switched into the wait state in t_4 . Therefore, both transactions idle.

Locking – Lost update problem 3/3



Transaction A	Time	Transaction B
READ t (S lock on t received)	t_1	-
-	t_2	READ t (S lock on t received)
WRITE t (X lock on t requested)	t_3	-
wait	t_4	WRITE t (X lock on t requested)
wait		wait
wait		wait

- We see that we solve one problem, however, another problem occurs. We call such a problem as **deadlock**.

Locking – Uncommitted dependency I 1/3



Transaction A	Time	Transaction B
-	t_1	WRITE t (X lock on t assigned)
READ t (S lock on t requested) wait	t_2	-
	t_3	COMMIT/ROLLBACK (X lock on t released)
repeat: READ t (S lock on t assigned)	t_4	

- The read of the transaction A in t_2 is not accepted due to the conflict with the X lock assigned to the transaction B in t_1 .
- The transaction A is switched into the wait state until the transaction B is finished by COMMIT or ROLLBACK and releases the exclusive lock on t .

Locking – Uncommitted dependency I 2/3



Transaction A	Time	Transaction B
-	t_1	WRITE t (X lock on t assigned)
READ t (S lock on t requested) wait	t_2	-
	t_3	COMMIT/ROLLBACK (X lock on t released)
repeat: READ t (S lock on t assigned)	t_4	

- After the release, A continues: it works with the valid value of t : if the transaction B proceeds COMMIT, then the transaction A works with the value set in t_1 , if B proceeds ROLLBACK then A works with the value set before t_1 .

Locking – Uncommitted dependency I 3/3



Transaction A	Time	Transaction B
-	t_1	WRITE t (X lock on t assigned)
READ t (S lock on t requested) wait	t_2	-
	t_3	COMMIT/ROLLBACK (X lock on t released)
repeat: READ t (S lock on t assigned)	t_4	

- Since the transaction A is not dependent on an uncommitted update of the transaction B , this problem is solved.

Locking – Uncommitted dependency II 1/3



Transaction A	Time	Transaction B
-	t_1	WRITE t (X lock on t assigned)
WRITE t (X lock on t requested) wait	t_2	-
	t_3	COMMIT/ROLLBACK (X lock on t released)
repeat: WRITE t (X lock on t assigned)	t_4	

- The write operation of the transaction A in t_2 is not accepted due to the conflict with the X lock of the transaction B .
- Therefore, the transaction A is switched into the wait state, until the transaction B is finished by COMMIT or ROLLBACK and releases the X lock on t .

Locking – Uncommitted dependency II 2/3



Transaction A	Time	Transaction B
-	t_1	WRITE t (X lock on t assigned)
WRITE t (X lock on t requested) wait	t_2	-
	t_3	COMMIT/ROLLBACK (X lock on t released)
repeat: WRITE t (X lock on t assigned)	t_4	

- The transaction A continues, it receives the X lock on t with the valid value: if B proceeds COMMIT, then A has the value set in t_1 , if B proceeds ROLLBACK, then A works with the value set before t_1 .

Locking – Uncommitted dependency II 3/3



Transaction A	Time	Transaction B
-	t_1	WRITE t (X lock on t assigned)
WRITE t (X lock on t requested) wait	t_2	-
	t_3	COMMIT/ROLLBACK (X lock on t released)
repeat: WRITE t (X lock on t assigned)	t_4	

- Since the transaction A is not dependent on an uncommitted update of the transaction B , this problem is solved.



Locking – Inconsistent analysis 1/3

$acc_1 = 30$	$acc_2 = 20$	$acc_3 = 50$
Transaction A	Time	Transaction B
READ acc_1 (S lock on acc_1 assigned) $suma = 30$	t_1	-
READ acc_2 (S lock on acc_2 assigned) $suma = 50$	t_2	-
-	t_3	READ acc_3 (S lock on acc_3 assigned)
-	t_4	WRITE $acc_3 = 60$ (X lock on acc_3 assigned)
-	t_5	READ acc_1 (S lock on acc_1 assigned)
...

Locking – Inconsistent analysis 2/3



$acc_1 = 30$	$acc_2 = 20$	$acc_3 = 50$
Transaction A	Time	Transaction B
...
-	t_6	WRITE $acc_1 = 20$ (X lock on acc_1 requested)
READ acc_3 (S lock on acc_3 requested)	t_7	wait
wait		wait

- The write of acc_1 by the transaction B in t_6 is not accepted due to the conflict of the X lock request with the S lock kept by the A transaction.
- The transaction B is switched into the wait state.



Locking – Inconsistent analysis 3/3

$acc_1 = 30$	$acc_2 = 20$	$acc_3 = 50$
Transaction A	Time	Transaction B
...
-	t_6	WRITE $acc_1 = 20$ (X lock on acc_1 requested)
READ acc_3 (S lock on acc_1 requested)	t_7	wait
wait		wait

- Similarly, the read of the transaction A in t_7 is not accepted since the conflict of the S lock request with the share lock of B . A is switched into the wait state.
- The original problem of inconsistent analysis is solved, however, **the deadlock occurs**.



It means the two-phase locking protocol can lead to a deadlock. In the following figure we see a general scheme of locking where the deadlock occurs.

Transaction A	Time	Transaction B
S lock on r_1 assigned	t_1	
-	t_2	S lock on r_2 assigned
X lock on r_2 requested	t_3	-
wait	t_4	X lock on r_1 requested
wait		wait

Remark: r_1 and r_2 are any database objects. They do not have to be tuples.



There are several techniques how to solve a deadlock:

1 Deadlock detection:

- 1 Time limits for a wait time,
- 2 Cycle detection in graph *Wait-For*.

2 Avoiding deadlock

Deadlock detection - time limits



- It is assumed that a transaction takes a **limited time**.
- When the transaction takes longer time, a system supposes that the deadlock occurs.

Deadlock detection - cycle detection



- A **cycle detection** is a more effective type of the deadlock detection. It creates a **graph of transactions** waiting for each other.
- The deadlock is solved by a selection of one deadlocked transaction for which the **ROLLBACK is processed** (i.e. all locks of the transaction are released).
- Remaining transactions in the deadlock can continue.
- A DBMS processes the cancelled transaction again or it returns a deadlock exception.

Avoiding deadlock 1/3



This strategy tries to avoid the deadlock by a modification of a locking protocol. There are two versions of locking protocol called **wait-die** and **wound-wait**:

- 1 Each transaction gets a timestamp – the time of the transaction start – which is unique.

Avoiding deadlock 2/3



- 2 If the transaction A requests a lock on a tuple, which is already locked by the transaction B , then:
 - 1 **Wait-die variant**: if A is older than B , then A is switched into the wait state; if A is younger than B , then A is canceled using ROLLBACK and runs again³.
 - 2 **Wound-wait variant**: if A is older than B , then the transaction B is cancelled using ROLLBACK and runs again⁴, if A is younger than B , then A is switched into the wait state;
- 3 If the transaction runs again the time-stamp remains the same.

³Transaction A die.

⁴Transaction B is wounded.

Avoiding deadlock 3/3



- The first part of the variant name describes a situation when A is older than B .
 - In the case of **wait-die**, the older transactions are switched into the wait state,
 - In the case of **wound-wait**, the younger transactions are switched into the wait state.
- It has been proved that the deadlock can not occur in the case of these protocols.
- A disadvantage of these protocols is relatively a high number of ROLLBACK operations.